

1 Supervised classification and regression

The goal of classification and regression is to find a function $f \in \mathcal{F}$, such that given an input \mathbf{x} , $f(\mathbf{x})$, “does what we want it to do”. We can define this as finding an f solving the following optimization

$$\min_{f \in \mathcal{F}} \mathbb{E}(l(f(\mathbf{x}), y))$$

where l is some loss function

Classification: $\mathbf{x} \in \mathbb{R}^p, y \in \{1, 2, \dots, k\}$

Regression: $\mathbf{x} \in \mathbb{R}^p, y \in \mathbb{R}$

1.1 Empirical risk minimization

In supervised learning, we have a training dataset, $(\mathbf{x}_i, y_i)_{i=1}^n$. A learning procedure is a mapping from the training dataset to some $f \in \mathcal{F}$. The standard paradigm is to try to minimize the empirical risk:

$$\hat{f} = \min_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n l(f(\mathbf{x}_i), y_i)$$

What are some exceptions to this? Tree based methods, CART learning procedure

2 Loss functions

2.1 Loss functions for binary classification

Let's define the output as $y \in \{-1, 1\}$

Zero-one loss

$$l(\hat{y}, y) = \mathbf{1}\{\hat{y} \neq y\}$$

Another way to write this:

$$l(\hat{y}, y) = \mathbf{1}\{\hat{y}y \leq 0\}$$

What's the problem with using this loss function? It is hard to optimize: one must use combinatorial search, and there is no gradient.

Surrogate losses

Let's say your function $f(\mathbf{x})$ actually outputs a number representing 'degree of certainty' (e.g. distance from a margin, or a probability). Then we can define a function like $f(u)$

$$l(\hat{y}, y) = f(yf(\mathbf{x}))$$

Examples:

$$\begin{aligned}f(u) &= \max(0, -u) \\f(u) &= \max(0, 1 - u) \\f(u) &= \log(1 + e^{-u})\end{aligned}$$

2.2 Loss function for regression

Output is some $y \in \mathbb{R}$. Define

$$l(\hat{y}, y) = (\hat{y} - y)^2$$

3 Statistical models

Principle of maximum likelihood is a way to derive common loss functions.

3.1 Logistic regression

A case of a generalized linear model:

$$g(P(y = 1 \mid \mathbf{x})) = \boldsymbol{\beta}^T \mathbf{x}$$

In the case of logistic regression, let

$$\begin{aligned}g(u) &= \log\left(\frac{u}{1-u}\right) \\g^{-1}(u) &= \frac{e^u}{1+e^u}\end{aligned}$$

Therefore the modeling assumption is:

$$\log \frac{P(y = 1 \mid \mathbf{x})}{P(y = -1 \mid \mathbf{x})} = \boldsymbol{\beta}^T \mathbf{x}$$

Let $g(\mathbf{x}) = \frac{e^{\boldsymbol{\beta}^T \mathbf{x}}}{1 + e^{\boldsymbol{\beta}^T \mathbf{x}}}$. So $1 - g(\mathbf{x}) = \frac{1}{1 + e^{\boldsymbol{\beta}^T \mathbf{x}}}$.

If $y = 1$, we have

$$R(\boldsymbol{\beta}) = \log\left(\frac{e^{\boldsymbol{\beta}^T \mathbf{x}}}{1 + e^{\boldsymbol{\beta}^T \mathbf{x}}}\right) = -\log(1 + e^{-\boldsymbol{\beta}^T \mathbf{x}}) = -\log(1 + e^{-y\boldsymbol{\beta}^T \mathbf{x}})$$

If $y = -1$, we have

$$R(\boldsymbol{\beta}) = \log\left(\frac{1}{1 + e^{\boldsymbol{\beta}^T \mathbf{x}}}\right) = -\log(1 + e^{\boldsymbol{\beta}^T \mathbf{x}}) = \log(1 + e^{-y\boldsymbol{\beta}^T \mathbf{x}})$$

Note this corresponds to the logistic loss function above.

3.2 Linear regression

Assumption: $y = \beta^T \mathbf{x} + \epsilon$, where $\epsilon \sim \mathcal{N}(0, \sigma^2)$.

$$\begin{aligned} R(\beta) &= \log \mathcal{N}(\beta^T \mathbf{x}, \sigma^2)(y) \\ &= \log \left[(2\pi\sigma^2)^{-1/2} \exp \left\{ -\frac{1}{2} (y - \beta^T \mathbf{x})^2 \right\} \right] \\ &= K - \frac{1}{2} (y - \beta^T \mathbf{x})^2 \end{aligned}$$

4 How to solve these things?

We want to solve

$$\min_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n l(f(x_i), y_i)$$

There have been decades of research on solving these kinds of problems. There is very rarely a closed form solution. Linear regression is a case where there is. The problem is

$$\min_{\beta} \frac{1}{n} \sum_{i=1}^n (y_i - \beta^T \mathbf{x}_i)^2$$

Let's switch to vector - matrix notation. Define

$$\begin{aligned} \mathbf{y} &= (y_1, \dots, y_n)^T \\ \mathbf{X} &= \text{design matrix} \end{aligned}$$

Then we rewrite

$$l(\beta) = \frac{1}{n} \|\mathbf{y} - \mathbf{X}\beta\|^2$$

Let's solve it! We ignore the $\frac{1}{n}$ term, because it is just a scaling.

$$\|\mathbf{y} - \mathbf{X}\beta\|^2 = \mathbf{y}^T \mathbf{y} - 2(\mathbf{X}\beta)^T \mathbf{y} + \beta^T \mathbf{X}^T \mathbf{X} \beta$$

Let's take the gradient w.r.t. β :

$$\nabla l = -2\mathbf{X}^T \mathbf{y} + 2\mathbf{X}^T \mathbf{X} \beta$$

Set it to zero

$$\begin{aligned} 0 &= -2\mathbf{X}^T \mathbf{y} + 2\mathbf{X}^T \mathbf{X} \beta \\ \mathbf{X}^T \mathbf{X} \beta &= \mathbf{X}^T \mathbf{y} \\ \beta &= (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \end{aligned}$$

This is a closed form solution, but it involves matrix multiplication. Also it requires $(\mathbf{X}^T \mathbf{X})^{-1}$ to exist.

4.1 Gradient descent

In general the function class \mathcal{F} is composed of some parametrized functions.

$$R(\beta) = \frac{1}{n} \sum_{i=1}^n l(f_{\beta}(x_i), y_i)$$

To minimize R , we can run gradient descent

$$\beta_{t+1} \leftarrow \beta_t - \eta \nabla R(\beta_t)$$

In higher dimensions, the directional derivative is

$$\begin{aligned} \nabla_{\mathbf{v}} R(\beta) &= \left. \frac{d}{dt} R(\beta + t\mathbf{v}) \right|_{t=0} \\ &= \mathbf{v}^T \nabla_{\mathbf{v}} R(\beta) \end{aligned}$$

Heuristic justification of gradient descent

$$\begin{aligned} R(\beta_{t+1}) &= R(\beta_t - \eta \nabla R(\beta_t)) \\ &\approx R(\beta_t) - \eta \nabla R(\beta_t)^T \nabla R(\beta_t) \\ &= R(\beta_t) - \eta \|\nabla R(\beta_t)\|^2 \end{aligned}$$

4.2 Stochastic gradient descent

The gradient of R is a sum over all training examples

$$\begin{aligned} \nabla R(\beta) &= \nabla \left(\frac{1}{n} \sum_i R_i(\beta) \right) \\ &= \frac{1}{n} \sum_{i=1}^n \nabla l(f_{\beta}(x_i), y_i) \end{aligned}$$

To evaluate one step of gradient descent, we have to look at all n examples and compute the loss. This is expensive, potentially very much so if the gradient is expensive to compute, as in a deep network. In stochastic gradient descent, we sample a single i , and calculate the gradient

$$\nabla R_i(\beta) = \nabla l(f_{\beta}(x_i), y_i)$$

The SGD update rule is:

$$\beta_{t+1} \leftarrow \beta_t - \eta \nabla R_i(\beta_t)$$

Why is it called stochastic gradient descent? Because of the random choice of i . By random sampling:

$$\mathbb{E}(\nabla R_i(\beta_t)) = \nabla R(\beta)$$

In general, SGD refers to a case where the update step is taken with a noisy gradient - i.e. one whose expectation is equal to the true gradient. The vast majority of the time, it refers this algorithm, in which an update is made to the parameters after visiting a single datapoint.

4.3 SGD for linear regression

$$R(\boldsymbol{\beta}) = \frac{1}{n} \sum_{i=1}^n \left(y_i - \boldsymbol{\beta}^T \mathbf{x}_i \right)^2$$

So $R_i = \left(y_i - \boldsymbol{\beta}^T \mathbf{x}_i \right)^2$. Let's calculate:

$$\nabla R_i(\boldsymbol{\beta}) = -2 \left(y_i - \boldsymbol{\beta}^T \mathbf{x}_i \right) \mathbf{x}_i$$

So what is the SGD update rule?

$$\boldsymbol{\beta}_{t+1} \leftarrow \boldsymbol{\beta}_t + \eta \left(y_i - \boldsymbol{\beta}^T \mathbf{x}_i \right) \mathbf{x}_i$$

This is the LMS algorithm! The predictor vector is added to the parameter, scaled by the residual.

4.4 SGD for classification

Let's use a linear classifier. Let's use the loss function $l(u) = \max(0, -u)$. Then

$$R(\boldsymbol{\beta}) = \frac{1}{n} \sum_{i=1}^n l \left(y_i \boldsymbol{\beta}^T \mathbf{x}_i \right)$$

So $R_i = \max(0, -y_i \boldsymbol{\beta}^T \mathbf{x}_i)$. What is

$$l'(u) = \begin{cases} -1 & \text{if } u < 0 \\ 0 & \text{if } u \geq 0 \end{cases}$$

Let's use the chain rule.

$$\nabla R_i(\boldsymbol{\beta}) = \begin{cases} -y_i \mathbf{x}_i & \text{if } y_i \boldsymbol{\beta}^T \mathbf{x}_i < 0 \\ 0 & \text{if } y_i \boldsymbol{\beta}^T \mathbf{x}_i \geq 0 \end{cases}$$

Does this look familiar? Let's say $y_i = 1$. If $\boldsymbol{\beta}^T \mathbf{x}_i < 0$, then this means the classifier put the example on the wrong side of the separating hyperplane. Then our update is

$$\boldsymbol{\beta}_{t+1} \leftarrow \boldsymbol{\beta}_t + \eta \mathbf{x}_i$$

This is the Perceptron update rule!

4.5 SGD for classification - logistic regression

Recall the logistic loss function:

$$R(\boldsymbol{\beta}) = \frac{1}{n} \sum_{i=1}^n \log \left(1 + e^{-y \boldsymbol{\beta}^T \mathbf{x}} \right)$$

Then calculate the gradient

$$\begin{aligned}\nabla R_i(\boldsymbol{\beta}) &= \frac{-y\mathbf{x}e^{-y\boldsymbol{\beta}^T\mathbf{x}}}{1+e^{-y\boldsymbol{\beta}^T\mathbf{x}}} \\ &= -y\mathbf{x} \left(1 - \text{logistic}(y\boldsymbol{\beta}^T\mathbf{x})\right)\end{aligned}$$

This depends on the fact that $\text{logistic}(-u) = 1 - \text{logistic}(u)$

Note that this is like a smoothed version of the perceptron rule: as the classifier is more confidently correct, the update becomes close to 0. As it is more confidently wrong, the update is close to $-y\mathbf{x}$, the perceptron rule.

Let's derive this another way (more akin to backpropagation). Consider the neural network

$$\begin{bmatrix} x_1 \\ \vdots \\ x_p \end{bmatrix} \rightarrow a \rightarrow p \rightarrow l$$

Where

$$\begin{aligned}a &= \boldsymbol{\beta}^T \mathbf{x} \\ p &= \frac{e^a}{1+e^a} \\ l &= -\mathbf{1}\{y=1\} \log p - \mathbf{1}\{y=-1\} \log(1-p)\end{aligned}$$

Let's calculate the gradient w.r.t. $\boldsymbol{\beta}$.

$$\nabla l = \frac{dl}{dp} \frac{dp}{da} (\nabla a)$$

Let's say $y = 1$:

$$\frac{dl}{dp} = -\frac{1}{p}$$

Then:

$$\frac{dp}{da} = \frac{d}{da} \left(\frac{e^a}{1+e^a} \right) = \frac{e^a + e^{2a} - e^{2a}}{(1+e^a)^2} = \frac{e^a}{(1+e^a)^2} = \left(\frac{e^a}{1+e^a} \right) \left(\frac{1}{1+e^a} \right) = p(1-p)$$

Also, $\nabla a = \mathbf{x}$.

So the full gradient becomes (for $y = 1$):

$$\nabla l = -\frac{1}{p} (p(1-p)) \mathbf{x} = -(1-p) \mathbf{x}$$

Now let's say $y = -1$. Then

$$\frac{dl}{dp} = \frac{1}{1-p}$$

So

$$\nabla l = \frac{1}{1-p} (p(1-p)) \mathbf{x} = p\mathbf{x}$$

5 Multiclass classification

How to generalize a binary classifier to a multi-class classifier? There are many ways. In deep learning, by far the most common is to use multinomial logistic regression / softmax activations with cross-entropy loss.

The assumption of multinomial logistic regression is that

$$P(y = j \mid \mathbf{x}) \propto \exp(\boldsymbol{\beta}_j^T \mathbf{x})$$

Or in other words

$$P(y = j \mid \mathbf{x}) = p_j = \frac{e^{\boldsymbol{\beta}_j^T \mathbf{x}}}{\sum_{l=1}^k e^{\boldsymbol{\beta}_l^T \mathbf{x}}}$$

This is known as the softmax function. The denominator ensures that the probabilities are normalized, i.e. that they sum to 1. This is the standard way to convert a vector of real-valued numbers into a probability distribution. We can think of this as a neural network:

$$\begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_p \end{bmatrix} \rightarrow \begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_k \end{bmatrix} \rightarrow \begin{bmatrix} p_1 \\ p_2 \\ \vdots \\ p_k \end{bmatrix}$$

Where $a_j = \boldsymbol{\beta}_j^T \mathbf{x}$. Note that it is invariant to constant shifts:

$$\frac{e^{C+\boldsymbol{\beta}_k^T \mathbf{x}}}{\sum_{j=1}^k e^{C+\boldsymbol{\beta}_j^T \mathbf{x}}} = \frac{e^C e^{\boldsymbol{\beta}_k^T \mathbf{x}}}{e^C \sum_{j=1}^k e^{\boldsymbol{\beta}_j^T \mathbf{x}}} = \frac{e^{\boldsymbol{\beta}_k^T \mathbf{x}}}{\sum_{j=1}^k e^{\boldsymbol{\beta}_j^T \mathbf{x}}}$$

Because of this, layer is overparametrized: we can arbitrarily set one of the activations to zero. We can set $a_k = 0$, and $a_j = \boldsymbol{\beta}_j^T \mathbf{x}$ for $j \in \{1, \dots, k-1\}$. Then we have a total of $p \times (k-1)$ parameters given by the vectors $\boldsymbol{\beta}_1, \dots, \boldsymbol{\beta}_{k-1}$.